

# Why Evolution Operates Near One Mutation per Genome per Generation

Ishanu Chatopadhyay<sup>1</sup><sup>a</sup>

This manuscript was compiled on February 13, 2026

Across organisms and viruses, the per-site mutation rate scales inversely with genome length, yielding an approximately constant number of mutations per genome per generation (Drake's rule). Existing explanations emphasize biochemical fidelity constraints or population-genetic error-threshold arguments that limit mutation from above. Here we provide a complementary, mechanism-independent account based on variation supply. Treating mutation as a blind local perturbation process, we ask which mutation rates maximize the expected rate of producing variants that are atypical relative to the mutation kernel itself, and therefore available for selective amplification rather than drift into mutational background. Using counting bounds for the fraction of such mutation-atypical outcomes within Hamming neighborhoods and combining them with an independent per-site mutation model, we obtain an explicit discovery-rate expression that is maximized when the per-genome mutation intensity  $T = n\mu$  is  $O(1)$ , with a peak near  $T \approx 1$ . Consequently, the optimal per-site rate scales as  $\mu^* = (1 + o(1))/n$ . These results derive Drake's rule from a generic supply–destruction tradeoff under local stochastic perturbations, complementing fidelity- and stability-based perspectives.

biological evolution | mutation rate | algorithmic complexity | Drake's rule

Mutation injects randomness; selection preserves and amplifies structure. Evolution therefore operates under a fundamental tension between entropy production and entropy suppression. A striking empirical regularity across diverse taxa is that the per-site mutation rate  $\mu$  scales approximately as  $1/n$ , where  $n$  is genome length, so the per-genome mutation rate  $U = n\mu$  is roughly constant. This phenomenon, observed from RNA viruses to eukaryotes, is known as Drake's rule (1–4).

Standard explanations of Drake's rule fall into two broad categories. One emphasizes biochemical constraints and tradeoffs in replication fidelity, proofreading, and energetic cost (2). The other emphasizes the population-genetic consequences of excessive mutation, most prominently Eigen's error-threshold framework (5), in which high mutation loads destabilize inherited information and limit the maintenance of adapted structure. These approaches provide important insight and useful upper bounds, but they primarily explain why mutation rates cannot be too large. They do not, by themselves, explain why evolution repeatedly operates near a particular scaling regime with  $U = O(1)$ .

Here we pursue a complementary route that is upstream of selection and does not invoke biochemical detail, explicit fitness landscapes, or equilibrium population genetics. We ask: given a blind local mutation mechanism, at what mutation load is the *supply rate* of statistically exceptional variants maximized? The point is not that evolution seeks exceptionality per se, but that selection can only act on what mutation supplies. If mutation predominantly produces outcomes that are typical under its own kernel, then repeated mutation pushes lineages toward the kernel's typical set, and accumulated organization is continually eroded into mutational background; we descend into noise devoid of structure making life impossible. Any sustained adaptive process therefore requires a nonzero supply of outcomes that are atypical relative to mutation alone, since these are the only candidates that can be preferentially retained rather than washed out.

Our starting observation is that mutation induces a natural statistical model for an offspring relative to its parent. For a parent genome  $x \in [q]^n$ , an  $m$ -mutation restricts the offspring  $y$  to the Hamming sphere  $H_m(x)$ , the set of sequences at Hamming distance  $m$  from  $x$ . Under the mutation-induced model that is uniform over  $H_m(x)$ , a typical draw requires  $\log |H_m(x)|$  bits to specify once the radius is fixed, and

$$|H_m(x)| = \binom{n}{m} (q-1)^m$$

Author affiliations: <sup>a</sup>University of Kentucky

IC conceived of research, implemented the algorithm, wrote the paper and procured support.

Authors have no competing interests.

<sup>2</sup>To whom correspondence should be addressed. E-mail: ishanu.ch@uky.edu

125 grows rapidly with  $m$ . This immediately implies a notion of  
 126 typicality: at a given radius, most outcomes resemble typical  
 127 draws from the mutation kernel, and only a small fraction are  
 128 exceptional relative to that kernel. In particular, the larger  
 129 the neighborhood, the thinner the tail of mutation-atypical  
 130 outcomes within it.

131 We formalize this atypicality or exceptionality relative  
 132 to the mutation-induced model on  $H_m(x)$  using randomness  
 133 deficiency (Definition S1 in the SI Appendix), which measures  
 134 how much rarer an outcome is than a typical draw from  
 135 the same model. Rarity here is not a proxy for fitness;  
 136 rather, being rare relative to the mutation kernel is what  
 137 makes an outcome statistically distinguishable from mutational  
 138 background and therefore *available*, in principle, for  
 139 selective amplification. A mutant is called a *net structure*  
 140 *discovery* if its deficiency exceeds a margin  $\Delta$  (Def. S1).  
 141 A two-part coding/counting argument (Theorem S1 and  
 142 Lemma S1) implies that, for typical parents, the fraction of  
 143  $\Delta$ -exceptional outcomes at radius  $m$  scales as  $2^{-\Delta}/|H_m(x)|$   
 144 up to polylogarithmic factors; this “typical” qualifier is  
 145 in a counting/combinatorial sense (formalized in the SI  
 146 Appendix) and is not an empirical claim that biological  
 147 genomes are algorithmically random. We then combine this  
 148 typical-case rarity with an independent per-site mutation  
 149 model  $M \sim \text{Binomial}(n, \mu)$  to define an expected discovery  
 150 rate  $\Phi(\mu)$  that averages the  $\Delta$ -exceptional probability over  
 151 mutation radii (Eq. S14). Optimizing  $\Phi(\mu)$  yields an interior  
 152 optimum at per-genome mutation intensity  $T = n\mu = O(1)$   
 153 (Theorem S2), and in particular,

$$\mu^* = \frac{1 + o(1)}{n}.$$

154 Thus Drake’s rule emerges as a mechanism-independent supply  
 155 optimum under blind local perturbation, complementing  
 156 fidelity- and stability-based perspectives such as the error-  
 157 threshold viewpoint.

## 158 Discussion

159 This work provides a mechanism-independent route to Drake’s  
 160 rule from the standpoint of variation supply. Under a  
 161 minimal model in which mutation acts as a local stochastic  
 162 perturbation on a discrete sequence space, we derived a  
 163 typical-case bound on the probability that an  $m$ -mutant is  
 164 statistically exceptional relative to its mutation neighborhood.  
 165 Combining this rarity scaling with an independent per-site  
 166 mutation model yields an explicit expression for the expected  
 167 rate at which mutation supplies such exceptional variants,  
 168 and shows that this rate is maximized when the expected  
 169 number of mutations per genome per generation is  $O(1)$ . In  
 170 the usual per-site parametrization, this corresponds to the  
 171 inverse-length scaling  $\mu^* = \Theta(1/n)$ .

172 A useful way to express the optimum is through the per-  
 173 genome mutation intensity

$$174 T := n\mu = \mathbb{E}[M], \quad M \sim \text{Binomial}(n, \mu),$$

175 which acts as an intensive “mutation temperature” controlling  
 176 the typical perturbation size. Figure 1 illustrates the scaling  
 177 directly. In the left panel,  $\Phi(\mu)$  plotted against the per-site  
 178 rate  $\mu$  exhibits a single interior maximum whose location  
 179 shifts left as  $n$  increases, consistent with  $\mu^* \propto 1/n$ . In  
 180 the right panel, plotting the same discovery rate against

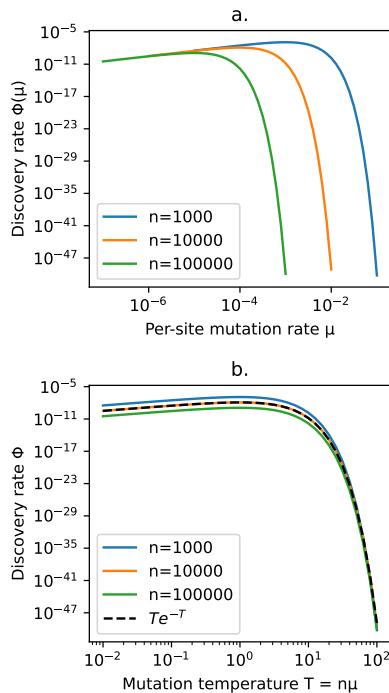


Fig. 1. Scaling and collapse of the discovery rate under independent per-site mutation. **a.** The expected discovery rate  $\Phi(\mu)$  as a function of the per-site mutation rate  $\mu$  for increasing genome lengths  $n$  (theory and Monte Carlo simulation). The location of the maximum shifts left as  $n$  increases, consistent with the inverse-length scaling  $\mu^* \propto 1/n$ . **b.** The same discovery rate plotted against the per-genome mutation intensity  $T = n\mu$ . Curves for different  $n$  collapse onto a common profile with a maximum near  $T \approx 1$ , matching the predicted  $Te^{-T}$  dependence up to a multiplicative constant. This collapse demonstrates that the optimal regime corresponds to a constant  $O(1)$  number of mutations per genome per generation.

181  $T = n\mu$  collapses curves for different  $n$  onto a common profile  
 182 with a maximum near  $T \approx 1$ , matching the predicted  $Te^{-T}$   
 183 dependence up to a multiplicative constant. The collapse  
 184 emphasizes that the scaling law is most naturally stated as a  
 185 constant optimal per-genome perturbation size.

186 Our notion of discovery is defined relative to the mutation-  
 187 induced model. We do not claim that the Hamming sphere  
 188 is the globally optimal statistical model for an offspring, only  
 189 that it is the model canonically induced by a local mutation  
 190 mechanism on sequence space. The substantive claim is  
 191 upstream: since mutation is the sole source of variation,  
 192 selection can only act on what mutation supplies. If offspring  
 193 were always typical draws from the mutation kernel, repeated  
 194 perturbation would concentrate lineages toward the kernel’s  
 195 typical set, eroding accumulated organization into mutational  
 196 background. Positive randomness deficiency provides a  
 197 quantitative, mechanism-independent way to express the  
 198 minimal sense in which a variant is not merely background  
 199 mutational noise.

200 The inverse-length scaling is robust to the choice of novelty  
 201 margin  $\Delta$ , provided  $\Delta$  does not grow with  $n$  fast enough to  
 202 dominate the one-step indexing cost  $c(1) = \Theta(\log n)$ . In  
 203 particular, any fixed  $\Delta$  or  $\Delta = o(\log n)$  yields the same  
 204  $\mu^* = \Theta(1/n)$  scaling, with  $\Delta$  affecting only the overall rate  
 205 through the factor  $2^{-\Delta}$ .

#1. we  
can  
probably  
add a  
figure  
pnale  
show-  
ing  
genomes  
are not  
very  
com-  
press-  
ible

249 **Why focus on discovery probability?** A natural objection to  
 250 the approach here could be to ask why evolution should be  
 251 viewed as optimizing the probability of structure discovery,  
 252 rather than a more direct quantity such as expected fitness  
 253 gain (2–4). We do not claim that natural selection explicitly  
 254 maximizes the objective we write down, nor do we attempt  
 255 to model fitness effects. Instead, the discovery probability  
 256 isolates an upstream constraint imposed by mutation itself  
 257 on the supply of potentially selectable variants. Mutation  
 258 generates variants blindly, without access to fitness *a priori*;  
 259 selection acts only after variants have been supplied. In this  
 260 setting it is well-posed to ask which per-site mutation rate  
 261 maximizes the expected rate at which mutation produces  
 262 outcomes that are atypical relative to the mutation kernel,  
 263 since such outcomes are precisely those that can, in principle,  
 264 be preferentially amplified rather than washed into the typical  
 265 mutational background.

266 It is also useful to situate this result relative to classical  
 267 “error threshold” arguments. Eigen-type frameworks constrain  
 268 mutation from above by requiring that inherited information  
 269 remain stable under copying errors (5). Our result is  
 270 different in kind: it identifies an interior optimum for the  
 271 upstream *supply* of mutation-atypical variants under blind  
 272 local perturbations, without invoking a fitness landscape  
 273 or equilibrium population genetics. The fact that both  
 274 perspectives emphasize the  $U = T = n\mu = O(1)$  regime  
 275 suggests that Drake’s rule may reflect the intersection of  
 276 two generic pressures: fidelity constraints that limit loss of  
 277 existing information and supply constraints that limit how  
 278 often mutation produces variants sufficiently atypical to be  
 279 preferentially amplified. In this sense, operating near one  
 280 mutation per genome per generation can be read as a regime  
 281 in which both retention and exploration remain nontrivial.

282 **Beyond biological specificity.** Although motivated by biological  
 283 evolution, the derivation depends only on three ingredients:  
 284 a discrete configuration space, a local stochastic perturbation  
 285 mechanism, and the requirement that adaptive progress  
 286 requires access to outcomes that are atypical relative to the  
 287 perturbation mechanism itself. Any system that explores  
 288 a high-dimensional space via local random perturbations  
 289 (a blind search unaware of downstream fitness functions)  
 290 faces an analogous tradeoff: perturbations that are too  
 291 small rarely generate statistically exceptional outcomes,  
 292 while perturbations that are too large overwhelm any low-  
 293 complexity deviation by pushing outcomes into exponentially  
 294 large neighborhoods in which the probability of exceptionality  
 295 is sharply suppressed. The resulting optimum selects an  $O(1)$   
 296 intensive perturbation scale, made visible here by the collapse  
 297 of discovery-rate curves under the parameter  $T$  in Fig. 1.

298 In summary, the combinatorial growth of mutation neighborhoods  
 299 suppresses the probability of mutation-atypical  
 300 outcomes at large radii, and optimizing the resulting discovery  
 301 rate under independent per-site mutation yields the observed  
 302 inverse scaling  $\mu^* \sim 1/n$ , equivalently an  $O(1)$  mutation load  
 303 per genome.

## 305 Materials and Methods

306 **Approach.** We model mutation as a blind local perturbation on  
 307  $[q]^n$ : conditional on  $m$  mutated sites, offspring are drawn uniformly  
 308 from the Hamming sphere  $H_m(x)$  around a parent  $x$  (formal setup

311 in the SI Appendix, Sec. 1). For any finite model class  $S \ni y$ , the  
 312 two-part code bound gives

$$313 K(y) \leq K(S) + \log |S| + O(1),$$

314 so describing  $y$  via  $S$  incurs an intrinsic indexing cost  $\log |S|$ .  
 315 Applying this with  $S = H_m(x)$  yields

$$316 K(y) \approx K(x) + \log |H_m(x)|, \quad |H_m(x)| = \binom{n}{m} (q-1)^m,$$

318 so the combinatorial growth of the mutation neighborhood determines  
 319 the typical description length of an  $m$ -mutant (details in  
 320 the SI Appendix, Eq. Eq. (S2)–Eq. Eq. (S3)).

321 We quantify mutation-atypicality using randomness deficiency  
 322 (Eq. Eq. (S5)), which measures how much rarer an outcome is than  
 323 a typical draw from the mutation-induced model. A mutant is  
 324 called a *net structure discovery* if its deficiency exceeds a margin  
 325  $\Delta$  (Def. S1). A counting argument (Thm. S1, Lem. S1) shows  
 326 that only a  $2^{-\Delta}/|H_m(x)|$  fraction of outcomes at radius  $m$  can be  
 327  $\Delta$ -exceptional, up to polylogarithmic factors.

328 We then embed this local rarity within an independent per-site  
 329 mutation model  $M \sim \text{Binomial}(n, \mu)$ , averaging over mutation  
 330 radii to obtain an expected discovery rate  $\Phi(\mu)$  (Cor. S1). Analyz-  
 331 ing the small- $\mu$  asymptotics of  $\Phi(\mu)$  (Lem. S2) and optimizing with  
 332 respect to  $\mu$  (Thm. S2) yields an interior optimum at  $n\mu = O(1)$   
 333 (Eq. Eq. (S20)).

334 **Code availability.** Code to generate the plots are available <https://github.com/zeroknowledgediscovery/drake>

335 **ACKNOWLEDGMENTS.** This work was supported in part by  
 336 the Defense Advanced Research Projects Agency (DARPA) under  
 337 the ARC program (HR0011-26-3-E016). Views, opinions, and  
 338 findings expressed are solely those of the authors.

1. JW Drake, A constant rate of spontaneous mutation in DNA-based microbes. *Proc. Natl. Acad. Sci.* **88**, 7160–7164 (1991).
2. M Lynch, Evolution of the mutation rate. *Trends Genet.* **26**, 345–352 (2010).
3. P Sniegowski, Y Raynes, Mutation rates: how low can you go? *Curr. Biol.* **23**, R147–R149 (2013).
4. R Hershberg, Mutation—the engine of evolution: studying mutation and its role in the evolution of bacteria. *Cold Spring Harb. Perspectives Biol.* **7**, a018077 (2015).
5. M Eigen, Selforganization of matter and the evolution of biological macromolecules. *Naturwissenschaften* **58**, 465–523 (1971).
6. M Li, PMB Vitányi, *An Introduction to Kolmogorov Complexity and Its Applications*. (Springer, 4 edition, (2019)).
7. P Gács, JT Tromp, PMB Vitányi, Algorithmic statistics. *IEEE Transactions on Inf. Theory* **47**, 2443–2463 (2001).

## Supplementary Methods

**Sequence Space, Mutation Geometry and Two-Part Codes.** Let  $[q]^n$  denote the set of length- $n$  strings over an alphabet of size  $q \geq 2$ . For  $x, y \in [q]^n$ , let  $d_H(x, y)$  denote the Hamming distance. For  $m \in \{0, 1, \dots, n\}$ , define the Hamming sphere

$$H_m(x) = \{y \in [q]^n : d_H(x, y) = m\}. \quad [\text{S1}]$$

with cardinality  $|H_m(x)| = \binom{n}{m}(q-1)^m$ , where  $m$  represents the number of mutated sites in one generation. Let  $K(\cdot)$  denote prefix-free Kolmogorov complexity with respect to a fixed universal Turing machine (6). If  $x \in S$  and  $S$  is a finite set, then the two-part code bound is as follows:

$$K(x) \leq K(S) + \log |S| + O(1), \quad K(x | S) \leq \log |S| + O(1). \quad [\text{S2}]$$

which expresses the optimality condition of two-part codes (6, 7): first describe the model  $S$ , then describe the index of  $x$  within  $S$ .

Next, we note that the mutation process induces a restricted model class. For fixed parent  $x$  and mutation radius  $m$ , the Hamming sphere  $H_m(x)$  is the set of sequences accessible by an  $m$ -mutation. Crucially,  $K(H_m(x) | x, m) = O(1)$ , which implies  $K(H_m(x)) = K(x) + O(\log n)$ , where the logarithmic term accounts for encoding  $(n, m, q)$ . This leads to the interpretation of *mutation indexing cost* as

$$c(m) = \log |H_m(x)| = \log \binom{n}{m} + m \log(q-1). \quad [\text{S3}]$$

Thus, random mutation induces the two-part code

$$K(y) \approx K(x) + c(m), \quad [\text{S4}]$$

for  $y$  uniformly drawn from  $H_m(x)$ . Finally, for a model  $S \ni y$ , the randomness deficiency (6, 7) is

$$\Delta_y(S) = K(S) + \log |S| - K(y). \quad [\text{S5}]$$

Thus, admissible models are restricted to mutation-accessible sets  $H_m(x)$ .

**Definition S1** (Net structure discovery). Fix  $\Delta \geq 1$ . A mutant  $y \in H_m(x)$  constitutes a net structure discovery if  $\Delta_y(H_m(x)) \geq \Delta$ , or equivalently,  $K(y) \leq K(x) + c(m) - \Delta$ .

Net structure discovery formalizes novelty as being exceptional relative to what mutation alone would typically produce.

**Theorem S1** (Universal upper bound). For any  $x \in [q]^n$ , any  $m$ , and any  $\Delta \geq 1$ ,

$$\Pr_{y \sim \text{Unif}(H_m(x))} [\Delta_y(H_m(x)) \geq \Delta] \leq O\left(\frac{2^{-\Delta}}{|H_m(x)|}\right). \quad [\text{S6}]$$

*Proof.* If  $\Delta_y(H_m(x)) \geq \Delta$ , then  $K(y) \leq K(H_m(x)) + \log |H_m(x)| - \Delta$ . By the Kraft inequality for prefix-free complexity, at most  $2^{T+1}$  strings satisfy  $K(y) \leq T$ . Substituting  $T = K(H_m(x)) + \log |H_m(x)| - \Delta$  and dividing by  $|H_m(x)|$  yields the bound.  $\square$

**Lemma S1** (Typical tightness for incompressible parents). Let  $x$  satisfy  $K(x) \geq n \log q - O(1)$ . Then for  $y$  drawn uniformly from  $H_m(x)$ ,

$$K(y | x, m) = \log |H_m(x)| \pm O(1), \quad [\text{S7}]$$

with probability  $1 - O(1/|H_m(x)|)$ . Consequently, up to polylogarithmic factors,

$$\Pr [\Delta_y(H_m(x)) \geq \Delta] = \Theta\left(\frac{2^{-\Delta}}{|H_m(x)|}\right). \quad [\text{S8}]$$

*Proof.* Fix  $x$  and  $m$ . There exists a computable bijection

$$\pi_{x,m} : \{1, \dots, |H_m(x)|\} \rightarrow H_m(x) \quad [\text{S9}]$$

such that given  $(x, m, i)$  one can compute  $y = \pi_{x,m}(i)$ . Hence

$$K(y | x, m) \leq \log |H_m(x)| + O(1) \quad [\text{S10}]$$

for all  $y \in H_m(x)$ . For the lower tail, define

$$A_\Delta = \{y \in H_m(x) : K(y | x, m) \leq \log |H_m(x)| - \Delta\}. \quad [\text{S11}]$$

Each  $y \in A_\Delta$  has a prefix-free description of length at most  $\log |H_m(x)| - \Delta$ , so  $|A_\Delta| \leq |H_m(x)|2^{-\Delta+O(1)}$ . Thus

$$\Pr[y \in A_\Delta] \leq 2^{-\Delta+O(1)}. \quad [\text{S12}]$$

Since  $x$  is incompressible, the two-part code through  $H_m(x)$  is optimal up to  $O(\log n)$  terms. Therefore

$$\Delta_y(H_m(x)) = \log |H_m(x)| - K(y | x, m) \pm O(\log n), \quad [\text{S13}]$$

which yields the stated asymptotic bound.  $\square$

**Corollary S1** (Expected discovery under binomial mutation). Let  $X \sim \text{Unif}([q]^n)$  be the parent sequence. Let  $M \sim \text{Binomial}(n, \mu)$  denote the number of mutated sites, and conditional on  $(X, M = m)$  let  $Y \sim \text{Unif}(H_m(X))$ . Define the expected discovery probability

$$\begin{aligned} \Phi(\mu) &:= \mathbb{P}(\Delta_Y(H_M(X)) \geq \Delta) \\ &= \sum_{m=1}^n \Pr(M = m) \Pr_{y \sim H_m(X)} [\Delta_y(H_m(X)) \geq \Delta]. \end{aligned} \quad [\text{S14}]$$

Then, for  $\Delta = o(\log n)$ ,

$$\Phi(\mu) \asymp 2^{-\Delta} \sum_{m=1}^n \Pr(M = m) \frac{1}{|H_m(X)|}. \quad [\text{S15}]$$

*Proof.* For  $X \sim \text{Unif}([q]^n)$ , the standard incompressibility bound

$$\Pr[K(X) \leq n \log q - t] \leq q^{-t} \quad [\text{S16}]$$

implies that  $X$  satisfies the incompressibility condition of Lemma S1 with overwhelming probability. On this event, Lemma S1 yields

$$\Pr_{y \sim H_m(X)} [\Delta_y(H_m(X)) \geq \Delta] \asymp \frac{2^{-\Delta}}{|H_m(X)|}. \quad [\text{S17}]$$

The contribution of the exceptional set  $\{X : K(X) \leq n \log q - t\}$  is at most  $q^{-t}$  and is negligible for  $t = \omega(1)$ . Substituting the tight bound into the definition of  $\Phi(\mu)$  gives the stated expression.  $\square$

Note that upon substituting the binomial mutation model and the sphere cardinality  $|H_m(X)| = \binom{n}{m}(q-1)^m$ , the combinatorial factor  $\binom{n}{m}$  appearing in  $\Pr(M = m)$  cancels with the identical factor in  $|H_m(X)|$ . Thus the multiplicity of mutation-location choices does not amplify discovery probability: it is already accounted for by the mutation process itself. What remains is the symbol-choice entropy  $(q-1)^m$  and the exponential suppression arising from the rarity of exceptional strings within each mutation-accessible model. We can therefore state the small- $\mu$  asymptotics of the discovery rate as follows.

**Lemma S2** (Small- $\mu$  asymptotics of discovery rate). For  $\mu = o(1)$ ,

$$\sum_{m=1}^n \left(\frac{\mu}{q-1}\right)^m = \frac{\mu}{q-1} + O(\mu^2), \quad (1-\mu)^n = e^{-n\mu+O(n\mu^2)}. \quad [\text{S18}]$$

Consequently,

$$\Phi(\mu) \asymp C 2^{-\Delta} \frac{\mu}{q-1} e^{-n\mu+O(n\mu^2)}, \quad [\text{S19}]$$

where  $C$  absorbs polylogarithmic factors.

*Proof.* The first expansion is the truncated geometric series. Substituting  $n \log(1-\mu) = -n\mu + O(n\mu^2)$  into the expression for  $\Phi(\mu)$  gives the claim.  $\square$

**Theorem S2** (Optimal mutation rate). In the regime  $\mu = o(1)$  with  $n\mu^2 = o(1)$ , the leading term  $\mu e^{-n\mu}$  is maximized at  $n\mu = 1$ . Consequently,

$$\mu^* = \frac{1 + o(1)}{n}. \quad [\text{S20}]$$

**Corollary S2** (Drake's rule). Optimizing information-theoretic structure discovery under random mutation yields an inverse scaling of per-site mutation rate with genome length.